Efficient R: practical solutions
Dr Colin Gillespie

0.1 Course R package

Before starting, we first install `drat`

```
install.packages("drat")
```

Then install the course R package via

```
drat::addRepo("rcourses")
install.packages("nclRefficient", type="source")
```

This R package contains copies of the practicals, solutions and data sets that we require. To load the package, use

```
library("nclRefficient")
```

Each practical corresponds to a chapter in the notes.

Practical 1

1. Reproduce the timing results in chapter 1 using the `benchmark` function from the `rbenchmark` package.

2. **Case study** In this example, we are going to investigate loading a large data frame. First, we'll generate a large matrix of random numbers and save it as a csv file:[1]

   ```
   N = 1e5
   m = as.data.frame(matrix(runif(N), ncol=1000))
   write.csv(m, file="example.csv", row.names=FALSE)
   ```

   We can read the file the back in again using `read.csv`

   ```
   dd = read.csv("example.csv")
   ```

   To get a baseline result, time the `read.csv` function call above.

   ```
   system.time(read.csv("example.csv"))
   ```

   ```
   ##    user  system elapsed
   ##   0.252   0.000   0.254
   ```

   We will now look ways of speeding up this step.

   (a) Use the `nrows` argument to set the number of rows that will be read from your file.[2]

   (b) Set `comment.char=""` to turn off interpretation of comments.

[1] If setting N=1e6 is too large for your machine, reduce it at bit. For example, N=50,000.

[2] Hint, use `nrow(m)` to determine how many rows are in your matrix.

(c) Explicitly define the classes of each column using `colClasses`
in `read.csv`, for example, if we have 1000 columns that all have
data type numeric, then:

```
read.csv(file="example.csv",
         colClasses=rep("numeric", 1000))
```

(d) Use the `save` and `load` functions:

```
save(m, file="example.RData")
load(file="example.RData")
```

Which of the above give the biggest speed-ups? Are there any
downsides to using these techniques? Do your results depend on
the number of columns or the number of rows?

```
## 1. Using RData files is the fastest - although
## you have to read the data in first.  Set
## colClasses also produces an good speed-up.

## 2. Setting colClasses R is no longer checking
## your data types.  If your data is changing - for
## example it's coming from the web or a database,
## this may be problem.

## 3. The results do depend on the number of
## columns, as this code demonstrates
N = c(1, 10, 100, 1000, 1000, 10000)
l = numeric(5)
for (i in seq_along(N)) {
    m = as.data.frame(matrix(runif(N[6]), ncol = N[i]))
    write.csv(m, file = "example.csv", row.names = FALSE)
    cc = rep("numeric", N[i])
    l[i] = system.time(read.csv("example.csv", colClasses = cc))[3]
}
l

## Notice that when we have a large number of
## columns, we get a slow down in reading in data
## set (even though we have specified the column
## classes). The reason for this slow down is that
## we are creating a data frame and each column has
## to be initialised with a particular class.
```

Practical 2

1. In this question, we'll compare matrices and data frames. Suppose
    we have a matrix, d_m

    ```
    ##For fast computers
    #d_m = matrix(1:1000000, ncol=1000)
    ##Slower computers
    d_m = matrix(1:10000, ncol=100)
    dim(d_m)

    ## [1] 100 100
    ```

    and a data frame d_df:

    ```
    d_df = as.data.frame(d_m)
    colnames(d_df) = paste("c", 1:ncol(d_df), sep="")
    ```

    (a) Using the following code, calculate the relative differences
        between selecting the first column/row of a data frame and
        matrix.

        ```
        benchmark(replications=1000,
                  d_m[1,], d_df[1,], d_m[,1], d_df[,1],
                  columns=c('test', 'elapsed', 'relative'))
        ```

        Can you explain the result? Try varying the number of replica-
        tions.

        ```
        ## Two things are going on here
        ## 1. The very large difference when selecting columns and rows (in data frames) is because the da

        ##2. Matrices are also more memory efficient:
        m = matrix(runif(1e4), ncol=1e4)
        d = data.frame(m)
        object.size(m)

        ## 80200 bytes

        object.size(d)

        ## 1120568 bytes
        ```

    (b) When selecting columns in a data frame, there are a few differ-
        ent methods. For example,

        ```
        d_df$c10
        d_df[,10]
        d_df[,"c10"]
        d_df[,colnames(d_df) == "c10"]
        ```

        Compare these four methods.

2. Consider the following piece of code:

```
a = NULL
for(i in 1:n)
  a = c(a, 2 * pi * sin(i))
```

This code calculates the values:

$$2\pi \sin(1), 2\pi \sin(2), 2\pi sin(3), \ldots, 2\pi sin(n)$$

and stores them in a vector. Two obvious ways of speeding up this code are:

- Pre-allocate the vector `a` for storing your results.
- Remove $2 \times \pi$ from the loop, i.e. at the end of the loop have the statement: `2*pi*a`.

Try the above techniques for speeding up the loop. Vary n and plot your results.

3. R is an interpreted language; this means that the interpreter executes the program source code directly, statement by statement. Therefore, every function call takes time.[3] Consider these three examples:

```
n = 1e6
## Example 1
I = 0
for(i in 1:n) {
  10
  I = I + 1
}
## Example 2
I = 0
for(i in 1:n){
  (((((((((((10))))))))))) 
  I = I + 1
}
## Example 3
I = 0
for(i in 1:n){
  ##This is a comment
  ##But it is still parsed
  ##So takes time
  ##But not a lot
  ##So don't worry!
  10
  I = I + 1
}
```

Using the `benchmark` function, time these three examples.

Practical 3: parallel programming

1. To begin, load the `parallel` package and determine how many cores you have

   ```
   library(parallel)
   detectCores()
   ```

2. Run the parallel `apply` example in the notes.

   - On your machine, what value of `N` do you need to use to make the parallel code run quicker than the standard serial version?

   - When I ran the benchmarks, I didn't include the `makeCluster` and `stopCluster` functions calls. Include these calls in your timings. How does this affect your benchmarks?

3. Run the dice game Monte-Carlo example in the notes. Vary the parameter `M`.[4]

   [4] Try setting `M=50` and varying `N`.

Solutions

Solutions are contained within this package:

```
library("nclRefficient")
vignette("solutions1", package="nclRefficient")
```