

Predictive Analytics: practical 2

Cross validation and the bootstrap

- Fit a linear regression model to the cars2010 data set with FE as the response, using EngDispl, NumCyl and NumGears as predictors.
- What is the training error rate (RMSE) for this model?
- Re-train your model using the validation set approach to estimate a test RMSE, make your validation set equivalent to a half of the full available data.
- How does this compare to the training error that we estimated above?
- Go through the same process using the different methods for estimating test error. That is leave one out and k -fold crossvalidation as well as bootstrapping.
- How do these estimates compare with the validation set approach?
- The object returned by train also contains timing information that can be accessed via the times component of the list. Which of the methods is fastest?
- Using k -fold cross validation to estimate test error investigate how the number of folds effects the resultant estimates and computation time.
- Experiment with adding terms to the model, transformations of the predictors and interactions say and use cross validation to estimate test error for each. What is the best model you can find? You can still use the validate and mark functions to look at how your models fair on the unseen data.

The data set can be loaded `data("FuelEconomy", package = "AppliedPredictiveModeling")`.

Hint: The training error can be found by taking the square root of the average square residuals. The `sqrt` and `resid` functions may be useful.

10-fold cross validation can be shown to be a good choice for almost any situation.

The `$` notation can be used pick a single list component.

Penalised regression

The diabetes data set in the lars package contains measurements of a number of predictors to model a response y , a measure of disease progression. There are other columns in the data set which contain interactions so we will extract just the predictors and the response. The data has already been normalized.

```
## load the data in
data(diabetes, package = "lars")
diabetesdata = cbind(diabetes$x, y = diabetes$y)
```

- Try fitting a lasso, ridge and elastic net model using all of the main effects, pairwise interactions and square terms from each of the predictors.¹
- Try to narrow in on the region of lowest RMSE for each model, don't forget about the `tuneGrid` argument to the `train` function. We can view the coefficients via

¹ Hint: see notes for shortcut on creating model formula. Also be aware that if the predictor is a factor a polynomial term doesn't make sense

`fraction = 0` is the same as the null model.

`y ~ (.) ^ 2` is short hand for a model that includes pairwise interactions for each predictor, so if we use this we should only need to add the square terms

```
coef = predict(mLASSO$finalModel,  
  mode = "fraction",  
  s = mLASSO$bestTune$fraction, # which ever fraction was chosen as best  
  type = "coefficients"  
)
```

- How many features have been chosen by the lasso and enet models?
- How do these models compare to principal components and partial least squares regression?